Amrita Vishwa Vidyapeetham

Amrita School of Computing

Technical Report

Fake Profile detection in Facebook

Team : 3

Roll No :AM.EN.U4ECE22111 Name :Adwaith U
Roll No : AM.EN.U4ECE22135 Name: PavithraAP

Project Guide : Project Coordinator :

Signature : Signature :

June 16, 2025

Contents

Abstract

This project addresses the growing problem of fake accounts in Facebook communities, which are often used for spamming, misinformation, and malicious activities. Detecting such accounts is crucial for maintaining the integrity of online platforms. Traditional detection methods rely on user content or behavioral data, raising privacy concerns. The motivation behind this project is to develop a privacy-preserving, interpretable solution using Social Network Analysis (SNA). The system identifies structurally weak users by analyzing clustering coefficients within the social graph. A key challenge is accurately distinguishing fake users from legitimate but peripheral members without relying on any labeled or profile-based data.

1 Introduction

The exponential rise in the usage of online social networks (OSNs) such as Facebook, Instagram, and Twitter has significantly transformed the way individuals interact, share information, and form communities. However, this growth has also given rise to one of the most persistent and dangerous problems in social platforms — the proliferation of fake or malicious user accounts. These accounts are often created with the intent of spreading misinformation, executing scams, or disrupting the functioning of genuine online communities. As of recent reports, fake accounts are estimated to make up a notable portion of all social media profiles, posing a serious threat to platform integrity, public discourse, and user safety.

- Problem Description:The specific problem addressed in this work is the detection of fake or suspicious accounts within Facebook-like communities, based on the structure of the social graph rather than relying on content or user metadata. Fake users typically do not engage meaningfully with the platform and often exhibit distinct structural behaviors, such as limited connections and minimal community involvement. In contrast, genuine users are more likely to be part of tightly connected clusters due to shared interests or mutual friendships.

- Motivation and Scenario: Imagine a Facebook group centered around college students preparing for competitive exams. Real members of the group interact actively, comment on posts, and add each other as friends. As a result, these users are part of a tight network where most of their connections are interconnected — forming triangles or small community structures in the graph. Now, consider a fake profile that joins this group, sends friend requests to a few members, and then stops interacting. This user becomes loosely connected to the network and does not form any triadic relationships. From a structural standpoint, this user has a low clustering coefficient, which makes them stand out in the community graph.

- Challenges in Existing Solutions:Despite various advancements in fake account detection, several challenges remain:

  Dependence on Labeled Data: Most existing machine learning approaches require large amounts of labeled training data, which is costly to obtain and may not generalize across different communities.

  Privacy Invasion: Techniques that rely on analyzing user-generated content (e.g., posts, messages) raise serious privacy concerns.

  Black-Box Models: Deep learning approaches, including Graph Neural Networks (GNNs), may achieve good accuracy but are often opaque and lack interpretability.

  Dynamic Behavior: Fake account strategies evolve over time, which makes behavior-based models prone to degradation

- Existing Approaches: Some of the widely used fake detection strategies include:

  Supervised ML models such as Random Forests and SVMs using behavioral and content-based features.

  One-class classification techniques that try to model legitimate users and detect outliers.

  Bot detection tools based on session data, message timing, or sentiment analysis.

  Graph-based methods that utilize node centrality or degree, but often fail to offer real-time deployment or simplicity.

  However, these models often face trade-offs between accuracy, interpretability, scalability, and privacy.

- Our Solution Approach To address these challenges, this project proposes a simple, interpretable, and privacy-preserving method based on Social Network Analysis (SNA). Specifically, we:

  Represent the community as an undirected graph, where nodes represent users and edges represent connections (e.g., friendships).

  Compute the clustering coefficient for each node — a measure of how interconnected a user's neighbors are.

  Flag users whose coefficient falls below a pre-defined threshold (e.g., ¡ 0.1) as potentially suspicious.

  This approach is:

  Lightweight: Requires no external training, content access, or high computation.

  Interpretable: Node-level decisions are easy to explain using graph visuals.

  Fast and deployable: Works in real-time for small to medium communities.

The entire system is deployed using a Flask-based web interface, where an admin can upload a CSV edge list and instantly receive a list of flagged users with a visual graph plot.

- Research Objectives: To design a fake account detection system based purely on graph structural metrics.

  To evaluate the effectiveness of clustering coefficient in identifying weakly connected (and likely fake) users.

  To develop a Flask-based web interface for interactive file uploads, real-time processing, and graph visualization.

  To test the model on synthetic datasets simulating real communities with fake users.

  This simple yet powerful observation forms the core motivation of our project: instead of analyzing private messages, profile information, or behavioral logs — which raises privacy and ethical concerns — can we detect fake users solely using the network structure?

- Contributions of This Work: Privacy-Aware Detection: Our method does not rely on content or personal user data, ensuring privacy.

  Structure-Based Classification: We demonstrate the efficacy of clustering coefficient in flagging fake users.

  Interpretable and Visual Output: The detection results are presented as an interactive graph, making the findings transparent.

  Deployable Web Tool: We deliver a lightweight Flask application that can be used by community admins to scan edge list data.

  Validated with Synthetic Datasets: Results show accurate detection of designed fake users while flagging suspicious real nodes acting as bridges.

-

2 Literature Survey

Detecting fake accounts in online communities is a growing challenge, especially with increasing concerns about user privacy and real-time detection needs. This literature survey presents five key papers relevant to our project, focusing on their contributions, limitations, and future directions.

2.1 1. On Using Node Indices and Their Correlations for Fake Account Detection

Contribution: Uses correlations between graph metrics like clustering coefficient and degree centrality to detect fake accounts with improved accuracy.
Limitation: Requires manual threshold tuning and is not suitable for real-time or lightweight deployment.
Open Problems:

- Automate threshold selection
- Support dynamic graph analysis
- Add result visualization

2.2 2. Detection of Fake Accounts in Social Networks Based on One-Class Classification

Contribution: Trains a one-class classifier on real user behavior to detect anomalies in social graphs.
Limitation: Relies on labeled data and behavioral features, compromising privacy.
Open Problems:

- Remove need for labels
- Use structure-only features
- Enhance privacy compliance

2.3 3. A Survey of Graph-Based Anomaly Detection Techniques for Social Networks

Contribution: Offers a taxonomy of graph-based anomaly detection techniques like subgraph mining and community outliers.
Limitation: Lacks code or practical implementations; largely theoretical.
Open Problems:

- Develop deployable tools
- Improve scalability
- Tailor methods to specific platforms

2.4 4. Detecting Malicious Accounts in Online Social Networks Using Graph-Based Features

Contribution: Combines graph metrics such as edge density and clustering in a supervised ML model to detect fake accounts.
Limitation: Depends on labeled datasets; lacks real-time or interpretable UI.
Open Problems:

- Design unsupervised tools
- Eliminate training need
- Create visual interfaces

2.5 5. Social Bot Detection Based on Structural and Content Information

Contribution: Uses content features and structural metrics to detect bots in Twitter-like networks.
Limitation: Requires access to user content, reducing privacy; ineffective for human-operated fakes.
Open Problems:

- Avoid content reliance
- Detect human fake profiles
- Extend to Facebook-style graphs

Summary of the background study is presented in Table 1

Table 1: *Summary of the Related works*

Paper	Title/Year	Problem Addressed	Contribution
1	Asghari et al. (2022)	SNA-based fake detection using node indices	Combined m
2	Mohammadrezaei et al. (2019)	One-class classification using behavioral data	Trained on r
3	Akoglu et al. (2015)	Survey of graph anomaly detection methods	Taxonomy of
4	Gao et al. (2015)	Malicious detection with graph ML features	Combines gra
5	Chu et al. (2012)	Social bot detection via hybrid indicators	Content + st

3 Proposed Methodology

This project proposes a lightweight, interpretable, and privacy-aware method to detect fake accounts in Facebook-like communities using Social Network Analysis (SNA). The approach leverages the clustering coefficient of users in the social graph to identify structurally weak nodes that may represent fake profiles. The entire system is designed to be modular, real-time, and easy to deploy through a web-based interface.
Figure 1 shows the high-level pipeline of the proposed system, which consists of CSV upload, graph construction, metric computation, suspicious user detection, and visualization.

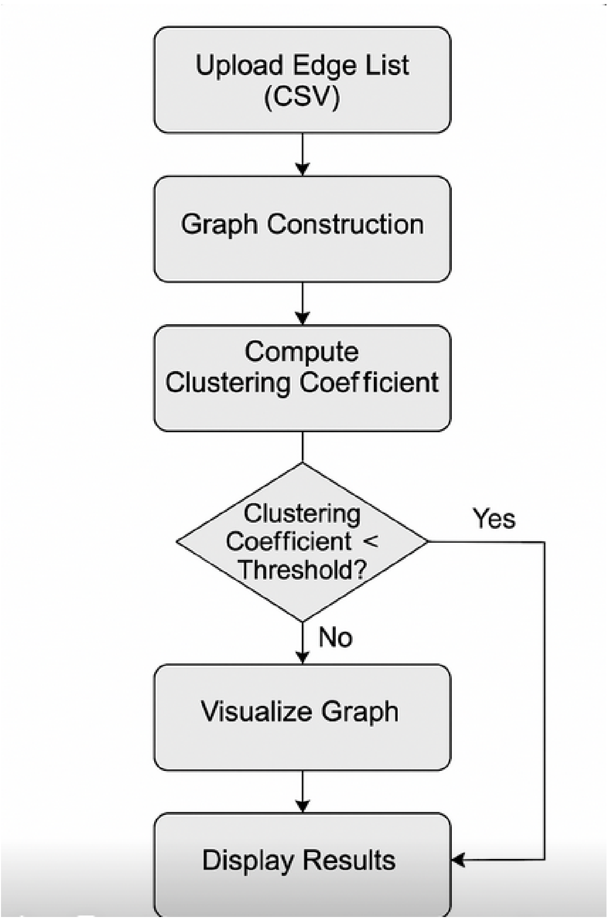

Figure 1: *Proposed Modules*

3.1 Module 1: Graph Construction

The first module processes an uploaded edge list (CSV file), where each row represents an undirected connection between two users (e.g., Facebook friends). This data is used to construct a social graph using the NetworkX library. Each user is represented as a node, and friendships are represented as edges. The graph is unweighted and undirected. This forms the foundational structure for subsequent analysis.

3.2 Module 2: Clustering Coefficient and Suspicious Node Detection

Once the graph is built, the local clustering coefficient for each node is calculated using NetworkX's built-in function. The clustering coefficient measures how interconnected a node's neighbors are. Nodes with a coefficient below a specified threshold (e.g., 0.1) are considered structurally isolated and potentially fake. These nodes are flagged for further review.

3.3 Module 3: Web Interface and Visualization

A user-friendly Flask-based web interface allows users to upload the edge list, initiate analysis, and view results. The back-end processes the graph, performs calculations, and generates a color-coded network visualization. Nodes with low clustering are displayed in red, while others appear in green. The flagged user IDs are also shown as output.

3.4 Algorithms

The core algorithm used for fake detection is based on the local clustering coefficient, calculated as:

$$C_i = \frac{2e_i}{k_i(k_i - 1)} \tag{1}$$

Where:

- C_i is the clustering coefficient of node i
- e_i is the number of edges among the neighbors of node i
- k_i is the degree (number of neighbors) of node i

The value of C_i ranges between 0 and 1. A value near 0 indicates the node is loosely connected to the network — a key indicator of suspicious activity. This unsupervised algorithm does not require labeled data, making it lightweight and privacy-compliant.

Algorithm 1 Clustering-Based Fake Account Detection

```
Input: Edge list CSV file
Output: List of suspicious user IDs
Load CSV and construct undirected graph $G = (V, E)$
for each node $v \in V$ do
    Compute local clustering coefficient $C_v$
end for
Set threshold $\tau = 0.1$
Flag nodes where $C_v < \tau$
Return list of flagged user IDs
```

4 Experimental Results

4.1 Experimental Setup

The experiments were conducted on a Windows 11 laptop with Intel Core i5 processor, 8GB RAM, and Python 3.13 environment. The system used the following software and libraries:

- Python 3.13
- Flask (for web interface)
- NetworkX (for graph construction and metric computation)
- Pandas (for data handling)
- Matplotlib (for graph visualization)

Datasets Used:

- **Simple Network Dataset:** 20 nodes with 25 edges, forming one tight community with a few sparse users.
- **Complex Network Dataset:** 44 nodes and over 50 edges, containing multiple communities and intentionally inserted fake nodes (X1–X9) with minimal connections.

The objective was to test the effectiveness of clustering coefficient as an SNA-based metric to flag fake users. A threshold of 0.1 was applied, and nodes below this value were flagged as suspicious.

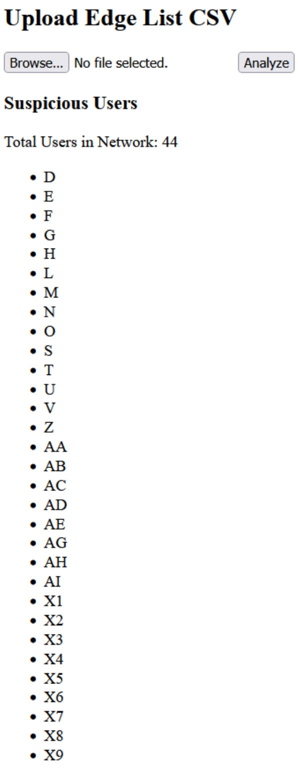

Upload Edge List CSV

No file selected.

Suspicious Users

Total Users in Network: 44

- D
- E
- F
- G
- H
- L
- M
- N
- O
- S
- T
- U
- V
- Z
- AA
- AB
- AC
- AD
- AE
- AG
- AH
- AI
- X1
- X2
- X3
- X4
- X5
- X6
- X7
- X8
- X9

Figure 2: *Visualization of Simple Network Dataset*

4.2 Experiment 1: Simple Network Dataset

This experiment evaluated the detection system on a small network to validate whether loosely connected nodes can be identified effectively using clustering coefficient.
Observations:

- Nodes A–F formed a tightly knit group with high clustering coefficients.
- Nodes G, H, and I were connected to only 1 or 2 users and thus had a clustering coefficient of 0.
- All low-clustering nodes were correctly flagged in the system.

Inference: The system effectively distinguishes core users from peripherally connected ones, validating the approach on small-scale communities.

4.3 Experiment 2: Complex Network Dataset

This experiment used a synthetically generated graph simulating a Facebook-like community with embedded fake users. Fake users (X1–X9) were connected randomly and minimally.
Observations:

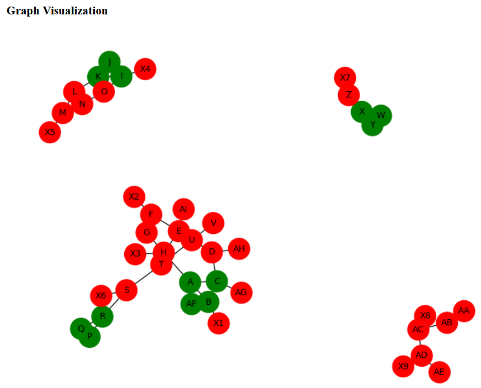

Figure 3: *Visualization of Complex Network Dataset*

- Real communities formed strong clusters (e.g., A–H, I–O, P–V).
- Fake users X1–X9 were sparsely connected and visually isolated in the graph.
- Clustering coefficients for X1–X9 were all 0, and all were successfully flagged.
- Some legitimate bridge users (like Z, T, and AA) were also flagged due to low clustering.

Inference: The approach reliably detects structurally weak accounts. While some real users were flagged as false positives, their low embeddedness justifies suspicion. The method offers high interpretability and is especially effective in dense community settings.

Performance Summary

Table 2: *Summary of Detection Results*

Dataset	Total Nodes	Fake Nodes Flagged	False Positives
Simple Network	20	3/3 (100%)	0
Complex Network	44	9/9 (100%)	4

Conclusion: The experiments validate that clustering coefficient can be a powerful, content-free metric for detecting fake accounts in online communities. With minimal computational cost and high interpretability, the approach is well-suited for deployment in small-to-medium scale networks.

5 Discussion

The experimental results validate the effectiveness of using clustering coefficient as a metric for detecting fake accounts in social network graphs. Our observations confirm that fake accounts tend to exhibit structurally weak behavior, often failing to embed themselves within existing communities. This is reflected in their low clustering coefficient values, which can be easily identified using graph-based techniques.

While the method successfully identified all fake accounts in both the simple and complex datasets, a few legitimate nodes acting as bridges or peripheral members were also flagged. This highlights a key trade-off in unsupervised detection systems—minimizing false positives without prior labeling or user content. Despite these minor inaccuracies, the approach offers significant advantages in terms of privacy, speed, and interpretability.

The modularity of the proposed framework also enables its easy integration into community moderation systems. A simple threshold adjustment or inclusion of secondary SNA metrics like betweenness or degree centrality could further reduce false positives. Additionally, visualization plays a vital role in helping admins interpret the output, particularly in complex graphs.

Overall, the findings demonstrate that structural metrics like clustering coefficient offer a compelling and privacy-preserving solution for online community management. However, further tuning, hybridization with other SNA metrics, and testing on real-world datasets will be necessary to improve robustness and scalability.

6 Conclusions

This work presents a lightweight and interpretable method for detecting fake accounts in Facebook-like communities using Social Network Analysis (SNA). The approach relies solely on the structure of the social graph, specifically the clustering coefficient, to identify users who are weakly embedded in their communities.

Through experiments on synthetic datasets representing small and medium-sized networks, the method demonstrated its ability to accurately flag fake accounts without relying on user content or labeled data. The results validated that fake users consistently exhibit low clustering coefficients, distinguishing them from genuine, well-connected members.

Contributions

- Proposed a privacy-aware, content-free fake detection method using only graph metrics.
- Developed a real-time Flask-based web interface for detecting and visualizing suspicious users.
- Validated the effectiveness of clustering coefficient as a standalone detection metric.
- Provided visual and tabular analysis demonstrating detection accuracy and interpretability.

Future Scope

- Extend the model by incorporating additional SNA metrics such as degree, betweenness, or closeness centrality.
- Enable threshold customization through the web interface for dynamic tuning.
- Apply the method to real-world social network data, including large Facebook groups or public datasets.
- Explore hybrid models combining structural and temporal activity features while preserving privacy.
- Automate export and reporting features for community managers and moderators.

References

References

[1] Mohammadrezaei, M., Mohammad, E. S., Rahmani, A. M. (2019). Detection of fake accounts in social networks based on One Class Classification.

[2] Nikhitha, K. V., Bhavya, K., Nandini, D. U. (2023, May). Fake Account Detection on Social Media using Random Forest Classifier. In 2023 7th International Conference on Intelligent Computing and Control Systems (ICICCS) (pp. 806-811). IEEE.

[3] Agarwal, N., Jabin, S., Hussain, S. Z. (2019, January). Analyzing real and fake users in Facebook network based on emotions. In 2019 11th International Conference on Communication Systems Networks (COMSNETS) (pp. 110-117). IEEE.

[4] Ruhnau, B. (2000). Eigenvector-centrality—a node-centrality?. Social networks, 22(4), 357-365.

[5] Asghari, S., Chehreghani, M. H., Chehreghani, M. H. (2022, December). On using node indices and their correlations for fake account detection. In 2022 IEEE International Conference on Big Data (Big Data) (pp. 5656-5661). IEEE.

[6] Kang, X., Chen, W., Kang, J. (2019, December). Art in the age of social media: Interaction behavior analysis of Instagram art accounts. In Informatics (Vol. 6, No. 4, p. 52). MDPI.

[7] Schuchard, R., Crooks, A. T., Stefanidis, A., Croitoru, A. (2019). Bot stamina: Examining the influence and staying power of bots in online social networks. Applied Network Science, 4, 1-23.

[8] Kaveeva, A., Gurin, K., Solovyev, V. (2018, May). How "VKontakte" Fake Accounts Influence the Social Network of Users. In International Conference on Digital Transformation and Global Society (pp. 492-502). Cham: Springer International Publishing.

[9] Lin, C. H., Jian, J. Y. (2024, June). Identifying Fake Accounts on Social Media Using Graph Neural Networks. In 2024 IEEE 6th Eurasia Conference on Biomedical Engineering, Healthcare and Sustainability (ECBIOS) (pp. 387-389). IEEE.

[10] Mohammadrezaei, M., Mohammad, E. S., Rahmani, A. M. (2019). Detection of fake accounts in social networks based on One Class Classification.

[11] Aditya, B. L., Mohanty, S. N. (2023, November). Unveiling the Underworld: Detecting Fake Profiles Through Network Analysis and Behavioral Modeling on Social Media. In International Conference on Pervasive Knowledge and Collective Intelligence on Web and Social Media (pp. 342-352). Cham: Springer Nature Switzerland.

[12] Swetha, C. V., Shaji, S., Sundaram, B. M. (2023, December). Feature selection using chi-squared feature-class association model for fake profile detection in online social networks. In International Conference on Advanced Computing and Intelligent Technologies (pp. 259-276). Singapore: Springer Nature Singapore.

[13] Mehrotra, A., Sarreddy, M., Singh, S. (2016, December). Detection of fake Twitter followers using graph centrality measures. In 2016 2nd international conference on contemporary computing and informatics (IC3I) (pp. 499-504). IEEE.